

Database Performance Benchmark

Explanation

This material explains the performance benchmarking script that measures PostgreSQL and MongoDB CRUD operations. It covers **what the code does**, **what the graphs mean**, and presents both **technical** and **non-technical** perspectives.

1. What the Code Does

1. Connects to databases

- PostgreSQL via `psycopg2`
- MongoDB via `pymongo`

2. Prepares test tables/collections

- PostgreSQL: creates `perf_test` table
- MongoDB: uses `perf_test` collection

3. Runs benchmarks for:

- **Single CRUD** (Create, Read, Update, Delete) → measures latency of a single operation
- **Bulk CRUD** (Create, Read, Update, Delete for 10,000 rows/documents) → measures latency and throughput

4. Calculates metrics

- **Average, minimum, and maximum time** per operation (in milliseconds)
- **Throughput** for bulk operations: how many rows/documents per second

5. Generates results

- **Rich table** in terminal for clear numerical comparison
- **Automatic conclusions** about which DB is faster for each operation
- **Matplotlib graphs**:
 - Top panel: average time per operation
 - Bottom panel: bulk throughput

2. Interpreting the Graphs

Top Panel: Average Time per Operation (ms)

- **Technical Explanation**

- Y-axis: time taken per operation in milliseconds
- X-axis: database operations
- Bars: skyblue = PostgreSQL, salmon = MongoDB
- Shows **latency for single and bulk operations**
- Shorter bars → faster completion → lower latency
- Critical for **interactive applications** where speed matters per query

- **Non-Technical Explanation**

- Think of this as a race:
 - Each bar shows how fast a database can do a job
 - Short bar = database completes the task faster
 - Single CRUD: small tasks (like adding a customer)
 - Bulk CRUD: bigger tasks (like adding 10,000 customers at once)
- You can see which database is faster at completing each task

Bottom Panel: Bulk Operations Throughput (rows/documents/sec)

- **Technical Explanation**

- Y-axis: number of rows/documents processed per second
- X-axis: bulk operations only
- Bars: same color scheme
- Shows **how well each database scales for large data operations**
- Taller bars = higher throughput → better performance for batch processing

- **Non-Technical Explanation**

- Imagine a cashier serving customers:
 - Taller bar = more people served per second
 - Shows which database handles a big crowd faster
 - Example: inserting 10,000 records
 - PostgreSQL = 40,000 rows/sec
 - MongoDB = 110,000 rows/sec → MongoDB completes the job faster

3. Key Takeaways

Metric	Technical Insight	Layman Insight
Single CRUD latency	Measures micro-performance; lower ms = faster response	Who finishes a small task faster
Bulk CRUD latency	Time taken for large batch operations	How long it takes to do a big job
Bulk throughput	Rows/documents per second; higher = better scaling	Who can handle more work at once
Overall trends	MongoDB generally faster for single inserts and high-volume writes	MongoDB is quicker for small tasks and big crowds; PostgreSQL may be competitive for some bulk reads

4. How to Explain the Results

1. For developers/technical audience:

- “The top panel shows per-operation latency (ms), highlighting MongoDB’s faster single inserts. The bottom panel shows bulk throughput; MongoDB outperforms PostgreSQL in bulk writes.”

2. For non-technical audience:

- “The top graph shows who can finish a task faster (shorter bars = faster). The bottom graph shows who can handle a crowd better (taller bars = more work done per second). MongoDB is faster both for small tasks and big jobs.”

5. Why This Matters

- **Latency (time per operation)** → affects user experience in real-time apps
- **Throughput (bulk operations)** → affects performance for analytics, migrations, and batch jobs
- Understanding both helps in **choosing the right database** for your workload.